

Quantum Game Theory: Equilibrium Search, Optimization, and Learning

John Bui

University of Maryland, Baltimore County
1000 Hilltop Circle Baltimore, MD 21250

johnlmbui@umbc.edu

Abstract

Quantum game theory extends classical strategic games by allowing players to apply unitary transformations to entangled quantum states. Within the Eisert–Wilkens–Lewenstein framework, equilibrium behavior is influenced by both entanglement and the admissible strategy space.

This work interprets quantum equilibrium search as a coupled optimization problem over continuous parameterized unitary strategy spaces. From this perspective, strategic adaptation resembles reinforcement learning and continuous policy optimization over probabilistic payoff landscapes.

We analyze the relationship between entanglement, equilibrium stability, admissible strategy classes, and optimization structure, while discussing limitations arising from nonconvexity, expanded strategy spaces, and physical noise.

1. Introduction

The prisoner’s dilemma is one of the canonical models in non-cooperative game theory that has become a central model for studying strategy selection and equilibrium stability [4]. Although mutual cooperation maximizes collective utility, rational players are driven toward mutual defection, producing a Nash equilibrium that is Pareto inefficient.

Quantum game theory extends classical strategic games into the quantum domain by replacing classical actions with operations on quantum states. In the Eisert–Wilkens–Lewenstein (EWL) formulation of the prisoner’s dilemma, players apply local unitary transformations to entangled qubits prior to measurement [3]. The resulting payoff distribution depends on both local unitary operations and the nonclassical correlations induced by entanglement.

Within quantum strategic games, entanglement changes how stable strategic outcomes emerge compared with classical game settings. Under maximal entanglement and re-

stricted strategy spaces, the classical equilibrium (D, D) is replaced by a cooperative quantum equilibrium (Q, Q) which is both Pareto optimal and strategically stable. However, subsequent work by Benjamin and Hayden showed that this equilibrium depends critically on assumptions regarding admissible strategy classes. Enlarging the strategy space introduces counter-strategies capable of destabilizing the EWL equilibrium.

From a computational perspective, quantum strategic games may also be interpreted as optimization problems over continuous unitary strategy spaces. Equilibrium search therefore resembles multi-agent optimization in which players iteratively adapt parameterized strategies to maximize expected payoff [5–7]. This interpretation connects quantum game theory to reinforcement learning and policy optimization.

The analysis investigates the relationship between entanglement, equilibrium stability, optimization, and statistical learning in quantum strategic games. More precisely, the structure of the strategy space, the dependence of equilibrium behavior on admissible unitary operations, and the interpretation of quantum equilibrium search as a continuous learning problem.

2. Background

2.1. Classical Prisoner’s Dilemma

The prisoner’s dilemma is a two-player non-zero-sum game involving the actions cooperate (C) and defect (D). The payoff ordering satisfies

$$T > R > P > S, \quad (1)$$

where T denotes the temptation payoff, R the reward for mutual cooperation, P the punishment for mutual defection, and S the sucker’s payoff.

A common numerical assignment is

$$T = 3, \quad R = 2, \quad P = 1, \quad S = 0. \quad (2)$$

Table 1. Payoff matrix for Bonnie and Clyde.

	Clyde cooperates	Clyde defects
Bonnie cooperates	(R, R)	(S, T)
Bonnie defects	(T, S)	(P, P)

Although (C, C) is Pareto superior, the Nash equilibrium is (D, D) because unilateral deviation from cooperation increases individual payoff.

2.2. Non-Zero-Sum Symmetric Games

The prisoner's dilemma is one example of a non-zero-sum symmetric game [9]. Another classical example is the game of chicken, in which players benefit from avoiding mutual conflict while attempting to avoid unilateral disadvantage. Unlike the prisoner's dilemma, chicken admits multiple Nash equilibria and emphasizes strategic coordination rather than dominant-strategy defection [6].

Table 2. Payoff matrix for the game of chicken played by players P_1 and P_2 .

	P_2 swerves	P_2 straight
P_1 swerves	$(0, 0)$	$(1, -1)$
P_1 straight	$(-1, 1)$	$(-1000, -1000)$

These examples illustrate how equilibrium structure may vary substantially across strategic environments, motivating the study of how quantum correlations alter equilibrium behavior in non-cooperative games.

2.3. Nash Equilibrium

A strategy profile s^* is a Nash equilibrium if no player can improve payoff by changing strategy alone:

$$u_i(s_i^*, s_{-i}^*) \geq u_i(s_i, s_{-i}^*) \quad (3)$$

for all admissible strategies s_i .

2.4. Quantum Computing Concepts

This analysis relies on standard quantum-mechanical principles including superposition, measurement, and entanglement.

Superposition. A qubit may exist in a linear combination of basis states:

$$|\psi\rangle = \alpha|0\rangle + \beta|1\rangle, \quad (4)$$

where

$$|\alpha|^2 + |\beta|^2 = 1. \quad (5)$$




Figure 1. Bloch sphere representation of a qubit state.

Measurement. Measurement collapses a quantum state probabilistically. The outcomes $|0\rangle$ and $|1\rangle$ occur with probabilities $|\alpha|^2$ and $|\beta|^2$ respectively.




Figure 2. Illustration of Schrödinger's cat thought experiment.

Entanglement. Two systems are entangled when their joint state cannot be decomposed into independent tensor products:

$$|\psi\rangle \neq |\psi_A\rangle \otimes |\psi_B\rangle. \quad (6)$$

Entanglement creates correlations that may not be reproducible through classical independent randomization.

3. Related Work

3.1. Meyer and Quantum Strategies

Meyer [8] introduced one of the earliest demonstrations of quantum strategic advantage. He showed that a player with access to quantum strategies may outperform an opponent restricted to classical strategies.

This work motivated the broader question of how strategic behavior changes when all players have access to quantum operations.

3.2. Eisert–Wilkens–Lewenstein Protocol

Eisert, Wilkens, and Lewenstein [3] proposed a quantum formulation of the prisoner’s dilemma. Their model applies entanglement to the initial game state and allows players to choose unitary operations as strategies.

Following the Eisert–Wilkens–Lewenstein (EWL) framework [3], maximal entanglement together with a restricted admissible strategy space may replace the classical equilibrium (D, D) with the cooperative quantum equilibrium (Q, Q) .

This result is significant because it suggests that entanglement may alter equilibrium structure itself rather than merely changing payoff magnitudes. From an optimization perspective, the EWL construction may be viewed as modifying the structure of the strategic search space through non-classical correlations.

Within the EWL framework [3, 11], each player selects a parameterized unitary strategy of the form

$$U(\theta, \phi),$$

producing expected payoff functions of the form

$$E[u_i(U_A, U_B; \gamma)].$$

3.3. Benjamin–Hayden Critique

Benjamin and Hayden [1] argued that the equilibrium structure of the EWL framework depends critically on restrictions placed on the admissible strategy space. If players are allowed to use larger classes of unitary operations, cooperative equilibria may lose stability, counter-strategies may emerge, and equilibrium structure may become significantly more complex.

Their critique reframes quantum game theory as a study of admissible strategy spaces and equilibrium stability rather than as a universal resolution of the prisoner’s dilemma.

From a learning and optimization perspective, this critique is vital because equilibrium stability becomes sensitive to representational assumptions, optimization constraints, and accessible policy classes. Consequently, quantum game theory may also be interpreted as a study of how strategic structure changes under different parameterizations of the admissible strategy space.

3.4. Learning and Optimization

Recent work in reinforcement learning and multi-agent learning [2, 11] motivates interpreting equilibrium search as optimization over continuous policy spaces. Because quantum strategies are parameterized by continuous unitary matrices, equilibrium computation resembles policy optimization in machine learning.

Although the original EWL framework was formulated primarily within quantum information theory, its structure closely resembles optimization over continuous policy spaces. Parameterized quantum strategies behave similarly to differentiable policy representations studied in reinforcement learning and continuous optimization.

4. Methods

4.1. Quantum Prisoner’s Dilemma

The following construction summarizes the standard EWL protocol [3]. The classical actions cooperate and defect are represented by the computational basis states

$$|C\rangle, |D\rangle. \quad (7)$$

The initial state is

$$|CC\rangle. \quad (8)$$

An entangling operator J is applied to create quantum correlations.

Each player then applies a local unitary strategy:

$$U_A \otimes U_B. \quad (9)$$

Finally, the entanglement is reversed using $J^\dagger$ and the state is measured.

The final state is

$$|\psi_f\rangle = J^\dagger(U_A \otimes U_B)J|CC\rangle. \quad (10)$$

After measurement, the final state produces probabilities

$$P_{CC}, P_{CD}, P_{DC}, P_{DD}. \quad (11)$$

Expected payoffs are computed as

$$E[u_A] = RP_{CC} + SP_{CD} + TP_{DC} + PP_{DD}, \quad (12)$$

and

$$E[u_B] = RP_{CC} + TP_{CD} + SP_{DC} + PP_{DD}. \quad (13)$$

4.2. Entanglement Operator

Let D denote the matrix representation of the classical defection operator. The entangling operator is commonly defined as

$$J(\gamma) = \exp\left(-\frac{i\gamma}{2}D \otimes D\right), \quad (14)$$

where

$$\gamma \in \left[0, \frac{\pi}{2}\right]. \quad (15)$$

Unlike classical mixed-strategy games, the resulting payoff functions depend continuously on unitary parameters and quantum correlations induced by entanglement. This continuous structure motivates interpreting equilibrium search as optimization over continuous parameterized unitary strategy spaces.

The parameter γ controls the degree of entanglement within the EWL framework. Specifically, $\gamma = 0$ corresponds to the classical game, whereas $\gamma = \pi/2$ corresponds to maximal entanglement.

4.3. Quantum Strategies

Players choose unitary strategies $U(\theta, \phi) \in SU(2)$ parameterized by angular variables (θ, ϕ) :

$$U(\theta, \phi) = \begin{pmatrix} e^{i\phi} \cos(\theta/2) & \sin(\theta/2) \\ -\sin(\theta/2) & e^{-i\phi} \cos(\theta/2) \end{pmatrix}. \quad (16)$$

Classical cooperation and defection appear as special cases:

$$C = I, \quad (17)$$

$$D = \begin{pmatrix} 0 & 1 \\ -1 & 0 \end{pmatrix}. \quad (18)$$

EWL also defines the quantum strategy

$$Q = \begin{pmatrix} i & 0 \\ 0 & -i \end{pmatrix}. \quad (19)$$

4.4. Equilibrium Search as Optimization

For fixed opponent strategy U_B , player A solves

$$U_A^* \in \arg \max_{U_A \in \mathcal{S}} E[u_A(U_A, U_B; \gamma)]. \quad (20)$$

Similarly, player B solves

$$U_B^* \in \arg \max_{U_B \in \mathcal{S}} E[u_B(U_A, U_B; \gamma)]. \quad (21)$$

A Nash equilibrium is therefore a fixed point of coupled optimization problems.

Unlike classical finite games, the quantum strategy space is continuous. Because admissible quantum strategies are parameterized continuously, equilibrium computation may be interpreted as a coupled nonconvex optimization problem over continuous unitary strategy spaces. Let each player parameterize a strategy by angular variables (θ, ϕ) and iteratively update these parameters through gradient ascent on the expected payoff functional:

$$(\theta_A, \phi_A) \leftarrow (\theta_A, \phi_A) + \eta \nabla E[u_A], \quad (22)$$

$$(\theta_B, \phi_B) \leftarrow (\theta_B, \phi_B) + \eta \nabla E[u_B]. \quad (23)$$

Here, $\eta > 0$ denotes a learning-rate parameter, while $\nabla E[u_i]$ represents the gradient of the expected payoff with respect to the player's strategy parameters. These updates are intended as conceptual analogies to policy-gradient learning rather than convergence-guaranteed optimization procedures.

This optimization perspective connects quantum equilibrium search to reinforcement learning and continuous policy optimization [2, 11].

From the perspective of statistical learning, expected payoffs may also be interpreted as stochastic objective functions generated by probabilistic measurement outcomes. Consequently, equilibrium search in quantum strategic games may be interpreted as optimization over continuous probabilistic policy spaces rather than equilibrium selection over finite discrete action spaces.

5. Discussion

Within the EWL formulation [3], admissible strategies are parameterized by continuous unitary transformations $U(\theta, \phi)$ acting on entangled quantum states. Consequently, equilibrium computation may be interpreted as optimization over continuous unitary strategy spaces rather than equilibrium search over finite discrete actions.

For fixed opponent strategy U_B , player A seeks to maximize the expected payoff functional

$$U_A^* \in \arg \max_{U_A \in \mathcal{S}} E[u_A(U_A, U_B; \gamma)], \quad (24)$$

with an analogous optimization problem for player B . This induces a coupled nonconvex optimization problem in which equilibrium structure is influenced by both the degree of entanglement and the structure of the admissible strategy space.

The surveyed literature suggests that entanglement may alter the strategic optimization landscape through nonclassical correlations in payoff generation.

Maximal entanglement may transform the classical Nash equilibrium (D, D) into the quantum equilibrium (Q, Q) , producing Pareto-superior outcomes [3]. However, the Benjamin–Hayden critique [1] demonstrates that these conclusions are highly sensitive to assumptions regarding the admissible strategy space.

This dependence is crucial because equilibrium stability becomes sensitive to parameterization choices, optimization constraints, accessible policy classes, and the representation of the payoff landscape. These observations parallel issues studied in statistical learning and continuous optimization, where convergence and stability are often governed by representation.

These observations suggest that quantum strategic games may be interpreted through the lens of statistical learning and continuous optimization, where equilibrium behavior

depends on both the payoff structure and parameterization of the admissible strategy space.

6. Conclusion

Quantum game theory extends classical strategic interaction into the quantum domain by enabling players to manipulate entangled quantum states through local unitary operations. Within the EWL framework, entanglement alters equilibrium structure and reveals cooperative quantum equilibria that do not emerge in the classical prisoner’s dilemma.

This work suggests that quantum equilibrium search may be interpreted as optimization over continuous parameterized unitary strategy spaces. Gradient-based interpretations of strategy adaptation suggest parallels with reinforcement learning and policy optimization.

Simultaneously, equilibrium stability is governed by the admissible strategy space. Restricted strategy classes enable cooperative equilibria, whereas expanded unitary spaces may introduce destabilizing counter-strategies.

As quantum computation and quantum machine learning continue to develop, quantum game-theoretic models may provide useful frameworks for studying strategic adaptation in computationally complex environments.

7. Limitations

Several limitations remain in current formulations of quantum game theory. First, equilibrium structure depends heavily on restrictions placed on the admissible strategy space. The cooperative equilibrium (Q, Q) identified in the Eisert–Wilkens–Lewenstein framework emerges only under particular classes of allowable unitary strategies. Enlarging the strategy space may introduce counter-strategies capable of destabilizing the equilibrium, as demonstrated by Benjamin and Hayden.

Second, most analyses of quantum strategic games remain mostly theoretical. Real quantum systems are affected by noise, decoherence, imperfect measurements, and operational errors. Consequently, theoretical equilibrium behavior derived under idealized assumptions may not persist in physically realizable quantum systems.

Third, equilibrium computation in continuous quantum strategy spaces produces generally nonlinear and nonconvex optimization landscapes. Although gradient-based interpretations resemble reinforcement learning and policy optimization, convergence guarantees for adaptive multi-agent learning dynamics remain vague in quantum strategic environments.

Moreover, the analysis does not include experimental validation and instead synthesizes findings from quantum game theory, optimization, and statistical learning literature. Future work may investigate numerical simulations, reinforcement learning approaches to equilibrium discov-

ery, and stability of equilibrium behavior under noise and expanded strategy classes.

Finally, it remains unclear whether quantum strategic advantages fundamentally produce new forms of strategic behavior or instead represent alternative formulations of correlated classical strategies. This question remains an active area of research in quantum game theory and quantum information science.

Acknowledgments

The author would like to thank Mohammadhossein Mohammadisiahroudi and Bedřich Sousedík for their mentorship and guidance, and Peter DeCrescenzo for his support through the UMBC Louis Stokes Alliance for Minority Participation (LSAMP). This work was supported, in part, by an Undergraduate Research Award from the UMBC Division of Undergraduate Academic Affairs. The author also acknowledges an introductory course in game theory taught by Jinglai Shen, which helped motivate this project.

References

- [1] Simon C. Benjamin and Patrick M. Hayden. Comment on “quantum games and quantum strategies”. *Physical Review Letters*, 87(6):069801, 2001. 3, 4
- [2] Lucian Buşoniu, Robert Babuška, and Bart De Schutter. A comprehensive survey of multiagent reinforcement learning. *IEEE Transactions on Systems, Man, and Cybernetics, Part C*, 38(2):156–172, 2008. 3, 4
- [3] Jens Eisert, Martin Wilkens, and Maciej Lewenstein. Quantum games and quantum strategies. *Physical Review Letters*, 83(15):3077–3080, 1999. 1, 3, 4
- [4] Thomas S. Ferguson. *A Course in Game Theory*. World Scientific, 2020. 1
- [5] Kazuki Ikeda and Shoto Aoki. Infinitely repeated quantum games and strategic efficiency. *Quantum Information Processing*, 20(12), 2021. 1
- [6] Richard Jankowski. Punishment in iterated chicken and prisoner’s dilemma games. *Rationality and Society*, 2(4):449–470, 1990. 2
- [7] Sau-Him Paul Lau and Vai-Lam Mui. Using turn taking to achieve intertemporal cooperation and symmetry in infinitely repeated 2×2 games. *Theory and Decision*, 72(2):167–188, 2012. 1
- [8] David A. Meyer. Quantum strategies. *Physical Review Letters*, 82(5):1052–1055, 1999. 2
- [9] John Nash. Non-cooperative games. *Annals of Mathematics*, 54(2):286–295, 1951. 2
- [10] Michael A. Nielsen and Isaac L. Chuang. *Quantum Computation and Quantum Information*. Cambridge University Press, 10th anniversary edition edition, 2010.
- [11] Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. MIT Press, 2nd edition, 2018. 3, 4